

Welcome to SENG 480B / CSC 485B / CSC 586B Self-Adaptive and Self-Managing Systems

Dr. Hausi A. Müller
Professor
Department of Computer Science
University of Victoria

<http://courses.seng.uvic.ca/courses/2013/summer/seng/480b>
<http://courses.seng.uvic.ca/courses/2013/summer/csc/485b>
<http://courses.seng.uvic.ca/courses/2013/summer/csc/586b>

Announcements

- Fri, June 28
 - Midterm in class
 - Prof. Venkatesh Srinivasan will administer the midterm
 - I will grade it ☺
- Midterm format
 - Closed books, closed notes, no phones or gadgets
 - Mostly essay type questions
- Midterm topics
 - Self-adaptive systems
 - Autonomic systems
 - MAPE-K loop
 - ACRA hierarchy
 - Symptoms
 - Policies
 - Feedback systems
 - Positive/negative/hybrid feedback
 - PID controllers
 - ULS systems
 - ULS characteristics
 - Wicked problems


Policy Types for Autonomic Computing

- Action policies
 - **If-then action rules** specify exactly what to do under the current condition.
 - Rational behaviour is compiled in by the designer
 - Basis for reflex agents
- Goal policies
 - Requires self-model, planning, conceptual knowledge representation
- Utility function policies
 - It chooses the actions to maximize its utility function
 - Finer distinction between desirability of different states than goals
 - Numerical characterization of state
 - Needs methods to carry out actions to optimize utility

Real autonomic systems embody a combination of policy types.

Action Policies


- A state S is a vector of attributes
- S can directly be measured by a sensor, or
- S can be inferred or synthesized from lower-level measurements
- Policy will directly or indirectly cause an action a
- Deterministic or probabilistic transition into a new state from S to a new state T



Kephart and Walsh: An Artificial Intelligence Perspective on Autonomic Computing Policies, POLICY 2004.

Action Policy Example

- RT: Response Time
 - if $(RT_{Gold} > 100 \text{ ms})$ increase CPU_{Gold} by 5%
 - if $(RT_{Silver} > 200 \text{ ms})$ increase CPU_{Silver} by 5%




Action Policy Examples

- For each machine, if idle session is greater than 20 minutes then terminate the session
- BitTorrent user processes initiated from IP address 141.223.2.15 should have lowest priority
 - if `(srcIPAddress == 141.223.2.15) && process-type == "bittorrent"` then priority is low
- Event
 - Total number of user logins is greater than 5 **and**
 - CPU load is greater than 90 **and**
 - Total number of processes running is greater than 35
- Action
 - Block any new user logins

Goal Policies

- Instead of specifying what to do in a current state, specify a single desired state or a set of desired states
- Any member of this target set is equally acceptable
- Cannot express fine distinctions in preference
- How to compute a set of actions that gets the system from current state S to a desired state T ?
- The system generates rational behaviour
- Potential conflicts
- How to resolve conflicts?




Kephart and Walsh: An Artificial Intelligence Perspective on Autonomic Computing Policies, POLICY 2004.

7

Goal Policy Example

- RT: Response Time
Gold: $RT_{\text{Gold}} \leq 100 \text{ ms}$
Silver: $RT_{\text{Silver}} \leq 200 \text{ ms}$



Kephart and Walsh: An Artificial Intelligence Perspective on Autonomic Computing Policies, POLICY 2004.

8

Utility-Function Policies

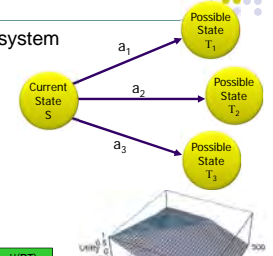
- An objective function to express the value of each possible state
- Generalizes goal policies
- Instead of desirable/undesirable we have a real-valued scalar desirability for each state
- More fine-grained and flexible specifications
- Goal functions often exhibit conflict; use utility function to resolve conflict
- Unambiguous, rational decision making

Kephart and Walsh: An Artificial Intelligence Perspective on Autonomic Computing Policies, POLICY 2004.

9

Utility Functions

- Map any possible state of a system to a scalar value
- Obtained from
 - Service Level Agreement
 - Preference elicitation
 - Simple templates
- Useful representation for high-level objectives
 - Value can be transformed to guide system behavior



Kephart and Walsh: An Artificial Intelligence Perspective on Autonomic Computing Policies, POLICY 2004.

10

Utility-Function Policy Example

- U : Utility function
 $U(RT_{\text{Gold}}, RT_{\text{Silver}}) = U_{\text{Gold}}(RT_{\text{Gold}}) + U_{\text{Silver}}(RT_{\text{Silver}})$



Kephart and Walsh: An Artificial Intelligence Perspective on Autonomic Computing Policies, POLICY 2004.

11

Policy Types for Autonomic Computing

- Real autonomic systems embody a combination of policy types
- In ACRA
 - Lower levels typically use action policies
 - For higher levels, goal or utility-function policies are more appropriate
- Unified framework is needed to support multiple policy types within a single autonomic component

Kephart and Walsh: An Artificial Intelligence Perspective on Autonomic Computing Policies, POLICY 2004.

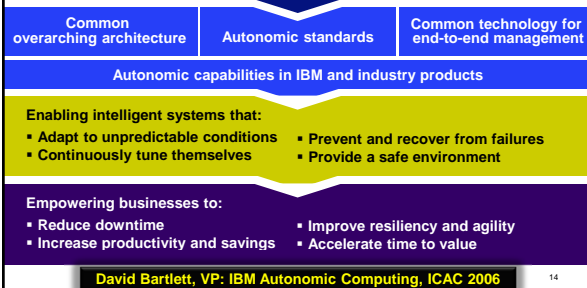
12

Policy Model and Terminology

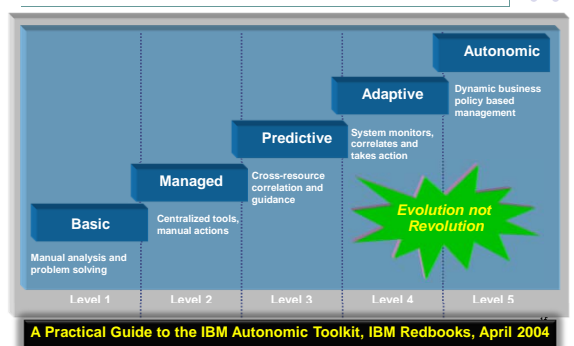
- **Policy**
 - A collection of policy alternatives.
- **Policy Alternative**
 - A collection of policy assertions.
- **Policy Assertion**
 - Represents an individual requirement, capability or other property of a behavior.
- **Policy Assertion Type**
 - Represents a class of policy assertions and implies a schema for the assertion and assertion-specific semantics.
- **Policy Assertion Parameter**
 - Qualifies the behavior indicated by a policy assertion.
- **Policy Vocabulary**
 - The set of all policy assertion types used in the policy.
- **Policy Expression**
 - An XML Infoset representation of a policy.
- **Policy Subject**
 - An entity (e.g., an endpoint, message, resource, interaction) with which a policy can be associated.
- **Policy Scope**
 - A collection of policy subjects to which a policy may apply.
- **Policy Attachment**
 - A mechanism for associating policy with one or more policy scopes.

IBM's Autonomic Computing Mission

Mission: Address IT complexity in heterogeneous environments through self-managing autonomic capabilities



Levels of Autonomic Computing Maturity



Levels of Autonomic Computing Maturity

	Basic Level 1	Managed Level 2	Predictive Level 3	Adaptive Level 4	Autonomic Level 5
Characteristics	Rely on system reports, product documentation, and manual actions to configure, optimize, heal and protect individual IT components	Management software in place to provide consolidation, facilitation and automation of IT tasks	Individual IT components and systems able to monitor, correlate and analyze the environment and recommend actions	IT components, individually and collectively, able to monitor, correlate, analyze and take action with minimal human intervention	Integrated IT components are collectively and dynamically managed by business rules and policies
Skills	Requires extensive, highly skilled IT staff	IT staff analyzes and takes actions	IT staff approves and initiates actions	IT staff manages performance against SLAs	IT staff focuses on enabling business needs
Benefits	Basic requirements addressed	Greater system awareness Improved productivity	Reduced dependency on deep skills Faster/better decision making	Balanced human/system interaction IT agility and resiliency	Business policy drives IT management Business agility and resiliency
	Manual				Autonomic

A Practical Guide to the IBM Autonomic Toolkit, IBM Redbooks, April 2004

Autonomic Computing Adoption Model

